

```
In [1]: import pandas as pd

df = pd.read_csv("../data/filteredvehpub.csv") # adjust path as need
df.head()
df.shape

## read files
```

```
Out[1]: (256115, 14)
```

```
In [2]: df.columns

## features and targets
```

```
Out[2]: Index(['HOUSEID', 'HHFAMINC', 'HHSIZE', 'LIF_CYC', 'CENSUS_R', 'URBA
N',
              'VEHID', 'DRVRCNT', 'MAKE', 'HH_RACE', 'WRKCOUNT', 'HOMEOWN',
              'URBANSIZE', 'VEHYEAR'],
              dtype='object')
```

```
In [3]: import numpy as np
invalid_values = [-9, -8, -7, -88, 99, "99", "XX", "xx", "XX ", "-88",
df = df.replace(invalid_values, np.nan)
df = df.dropna()

## Remove NaN
```

```
In [4]: #df["VEHTYPE"] = df["VEHTYPE"].replace([5, 6], np.nan)
#df = df.dropna()
```

```
In [5]: counts = df["MAKE"].value_counts()
counts.describe()

# Get percentile of the data
```

```
Out[5]: count      53.000000
mean      4607.245283
std       8163.852477
min        58.000000
25%       286.000000
50%      1580.000000
75%      4639.000000
max      34768.000000
Name: count, dtype: float64
```

```
In [6]: df["MAKE"] = df["MAKE"].astype(str) # convert everything to string f
df["MAKE"] = df["MAKE"].str.strip() # remove whitespace
df["MAKE"] = df["MAKE"].astype(int) # convert to integer

## Did this because the original data has different variable types, wa
```

```
counts = df["MAKE"].value_counts()
rare_makes = counts[counts < 1743].index
## The number of 1000 can be found in the report

df["MAKE"] = df["MAKE"].where(~df["MAKE"].isin(rare_makes), 98)
df["MAKE"].value_counts()
## Put them all into to other
```

```
Out[6]: MAKE
12      34768
49      33623
20      31188
37      23829
98      16040
7       12545
35      11684
2        7293
48      6932
23      6896
55      6765
59      4967
18      4852
6       4703
63      4639
30      4608
34      4536
41      4335
42      3642
72      3064
19      2664
22      2659
54      2536
51      1831
14      1798
13      1787
Name: count, dtype: int64
```

```
In [7]: #df.to_csv("processed_data.csv", index=False)
```

```
In [8]: df.columns.tolist()
```

```
Out[8]: ['HOUSEID',
        'HHFAMINC',
        'HHSIZE',
        'LIF_CYC',
        'CENSUS_R',
        'URBAN',
        'VEHID',
        'DRVRCNT',
        'MAKE',
        'HH_RACE',
        'WRKCOUNT',
        'HOMEOWN',
        'URBANSIZE',
        'VEHYEAR']
```

```
In [9]: feature_cols = [
        "HHSIZE",
        "HHFAMINC",
        "LIF_CYC",
        "CENSUS_R",
        "HH_RACE",
        "HOMEOWN",
        "WRKCOUNT",
        "URBAN",
        "URBANSIZE",
        "DRVRCNT"
    ]
    target_col = "MAKE"
    X = df[feature_cols]
    y = df[target_col]

    # Declare features and targets
```

```
In [10]: from sklearn.compose import ColumnTransformer
        from sklearn.preprocessing import OneHotEncoder
        from sklearn.preprocessing import StandardScaler

        numeric_features = ["HHSIZE", "HHFAMINC", "WRKCOUNT", "DRVRCNT"]
        categorical_features = ["LIF_CYC", "CENSUS_R", "HH_RACE", "HOMEOWN", "VEHYEAR"]

        preprocess = ColumnTransformer(
            transformers=[
                ("num", StandardScaler(), numeric_features),
                ("cat", OneHotEncoder(handle_unknown="ignore"), categorical_features)
            ]
        )

        ## Separate numeric features with categorical features
        ## Then use standardScaler() for num and One-hot encoding for cat
```

```
In [11]: from sklearn.pipeline import Pipeline
```

```

from sklearn.ensemble import RandomForestClassifier


rf = RandomForestClassifier(
    n_estimators=300,
    random_state=42,
    n_jobs=-1,          # use all CPU cores
    min_samples_split=4,
    min_samples_leaf=2
)

pipe = Pipeline([
    ("preprocess", preprocess),  # your ColumnTransformer from Step 7
    ("rf", rf)
])

pipe

## Construct the tree and the pipeline

```



```

In [12]: ## Cross Validation Accuracy

from sklearn.model_selection import StratifiedKFold
from sklearn.metrics import accuracy_score
from tqdm import tqdm
import numpy as np

cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

scores = []

for train_idx, test_idx in tqdm(cv.split(X, y), total=5):
    X_train_cv, X_test_cv = X.iloc[train_idx], X.iloc[test_idx]
    y_train_cv, y_test_cv = y.iloc[train_idx], y.iloc[test_idx]

    # Fit pipeline
    pipe.fit(X_train_cv, y_train_cv)

    # Evaluate accuracy

```

```
y_pred_cv = pipe.predict(X_test_cv)
scores.append(accuracy_score(y_test_cv, y_pred_cv))

print("Cross-validated Accuracy:", np.mean(scores), "+/-", np.std(scores))
```

100%|██████████| 5/5 [03:30<00:00, 4
2.00s/it]


Cross-validated Accuracy: 0.17221029061266055 +/- 0.0012839465344754422

```
In [13]: ## Fit data into 80% test and 20% train, to train the final model

from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split( X, y, test_size=0
```

```
In [14]: pipe.fit(X_train, y_train)
```



```
In [35]: ## Accuracy of the 80-20 test (seems like people do it for confusion m

from sklearn.metrics import accuracy_score

y_pred = pipe.predict(X_test)

test_acc = accuracy_score(y_test, y_pred)
print("Test Top-1 Accuracy:", test_acc)
```

Test Top-1 Accuracy: 0.17050596883510452

```
In [36]: ## Confusion Matrix

from sklearn.metrics import confusion_matrix
import matplotlib.pyplot as plt
import numpy as np

cm = confusion_matrix(y_test, y_pred)

labels = sorted(y.unique())
```

```


fig, ax = plt.subplots(figsize=(10, 10))
im = ax.imshow(cm, cmap="Blues")
plt.colorbar(im)

ax.set_xticks(np.arange(len(labels)))
ax.set_yticks(np.arange(len(labels)))
ax.set_xticklabels(labels, rotation=45, ha="right")
ax.set_yticklabels(labels)

ax.set_xlabel("Predicted")
ax.set_ylabel("True")
ax.set_title("Confusion Matrix (Top-1 Predictions)")

plt.tight_layout()
plt.show()

```



In []:

In []: